

# Crypto Volatility Spike Detection

## Hard Dataset – Model Comparison Report

Generated: 2025-11-19 11:55:34

HARD Train file: features\_hard\_train\_20251119\_114901.parquet

HARD Test file : features\_hard\_test\_20251119\_114901.parquet

### 1. Dataset Summary (Hard, Time-Based Split)

---

Train rows : 788

Test rows : 198

Number of features : 42

Train spike rate : 2.54%

Test spike rate : 15.15%

Train time range : 2025-11-07 22:09:00.382094+00:00 → 2025-11-08 15:20:17.136652+00:00

Test time range : 2025-11-08 15:20:17.433996+00:00 → 2025-11-08 15:31:16.274429+00:00

Construction notes:

- Label is a FUTURE volatility spike (proxy computed from future\_volatility window).
- Train/test split is strictly time-based (no shuffling) to simulate live deployment.
- All volatility-like and future-looking columns are removed from features to avoid leakage.
- Simple median imputation is used for missing numeric values.

### 2. Modeling & Evaluation Setup

---

Four models are trained on the HARD TRAIN split and evaluated on the HARD TEST split:

- Logistic Regression
- Random Forest
- Gradient Boosting
- XGBoost

All models:

- Use the same engineered numeric feature set (after leakage removal).
- Are trained only on past information (no access to future volatility).
- Are evaluated on the hard test set with ~15% spike rate.

Metrics reported:

- Accuracy – overall correctness (dominated by the majority class).
- AUC – ranking quality for spike vs non-spike.
- Precision – fraction of predicted spikes that are real spikes.
- Recall – fraction of true spikes that the model catches.
- F1 – harmonic mean of precision and recall (primary selection metric).

### 3. Model Comparison (Hard Test Set)

---

Model	Thr	Acc	AUC	Prec	Rec	F1
★ random_forest_hard_20251119_	0.260	0.848	0.835	0.000	0.000	0.000
gradient_boosting_hard_20251	0.010	0.838	0.654	0.000	0.000	0.000
xgboost_hard_20251119_114951	0.260	0.848	0.564	0.000	0.000	0.000
logistic_regression_hard_202	0.880	0.848	0.500	0.000	0.000	0.000

### 4. Best Performing Model (Winner)

---

Winner model (by F1, then AUC):

- Model name : random\_forest\_hard\_20251119\_114950
- Accuracy : 0.848
- AUC : 0.835
- Precision : 0.000
- Recall : 0.000
- F1 score : 0.000
- Decision threshold: 0.260

Interpretation:

- The winner is selected to maximize F1, which balances precision and recall.
- A higher AUC indicates better ranking quality, even if classification at a fixed threshold is conservative.
- In this hard, time-based setup with rare spikes, even modest recall